

An Introduction to the Database Management Systems

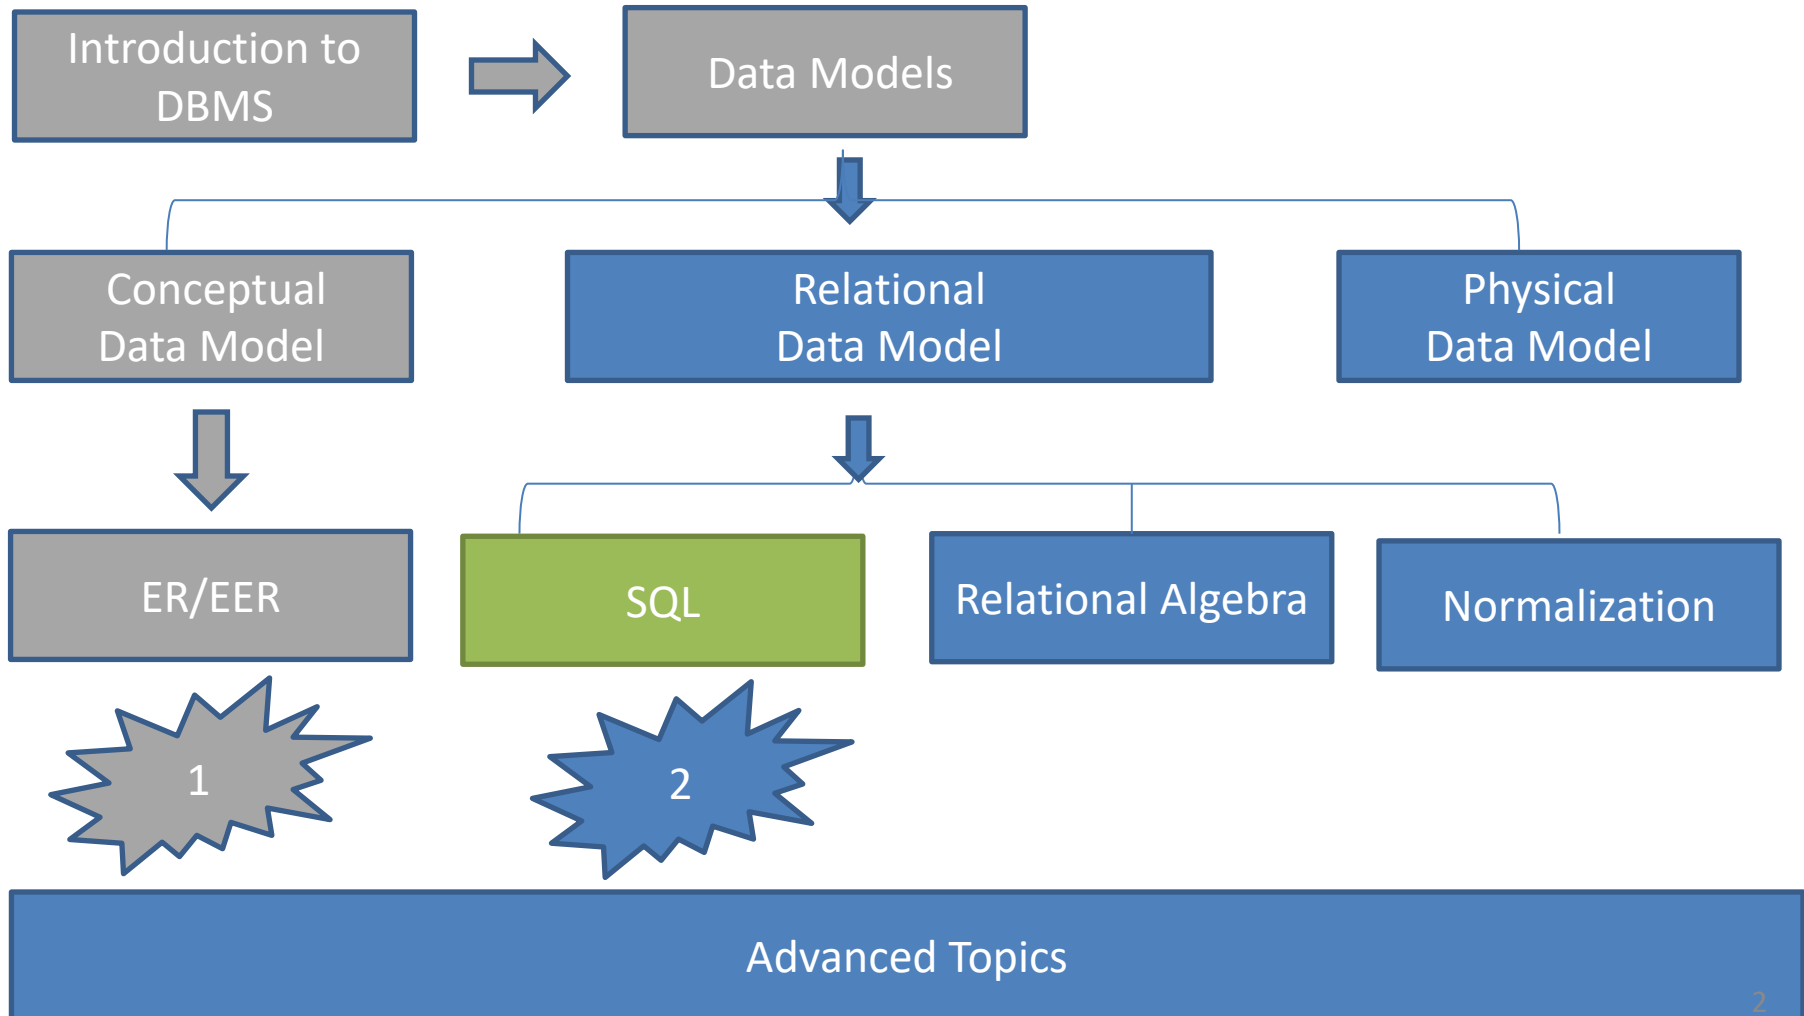
By
Hossein Rahmani

Slides originally by Book(s) Resources




Road Map

(Might change!)



SQL Advanced

- More Complex SQL Retrieval Queries 
- Specifying Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Change Statements in SQL

More Complex SQL Retrieval Queries

- Additional features allow users to specify more complex retrievals from database:
 - Nested queries, joined tables, outer joins, aggregate functions, and grouping

Comparisons Involving NULL and Three-Valued Logic

- Meanings of NULL
 - Unknown value
 - Unavailable or withheld value
 - Not applicable attribute
- Each individual NULL value considered to be different from every other NULL value
- SQL uses a three-valued logic:
 - TRUE, FALSE, and UNKNOWN

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

Table 5.1 Logical Connectives in Three-Valued Logic

(a)	AND	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	OR	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	NOT			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

Comparisons Involving NULL and Three-Valued Logic (cont'd.)

- SQL allows queries that check whether an attribute value is NULL
 - IS or IS NOT NULL

Query 18. Retrieve the names of all employees who do not have supervisors.

```
Q18:  SELECT  Fname, Lname
      FROM    EMPLOYEE
      WHERE   Super_ssn IS NULL;
```

Nested Queries, Tuples, and Set/Multiset Comparisons

- **Nested queries**
 - Complete select-from-where blocks within WHERE clause of another query
 - **Outer query**
- Comparison operator **IN**
 - Compares value v with a set (or multiset) of values V
 - Evaluates to **TRUE** if v is one of the elements in V

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

Figure 9.1
The ER conceptual schema diagram for the COMPANY database.

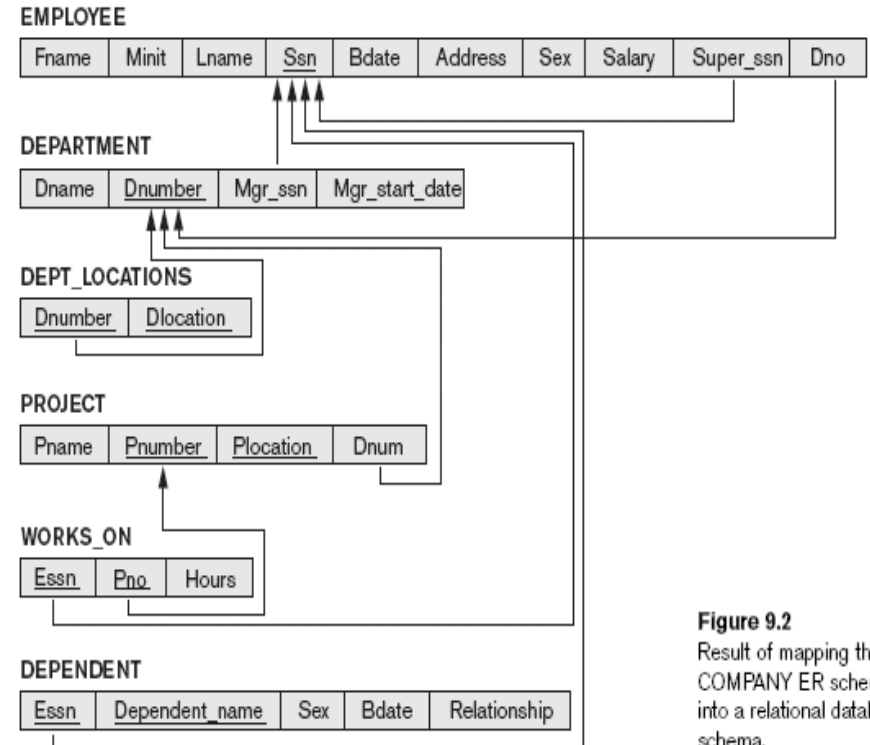
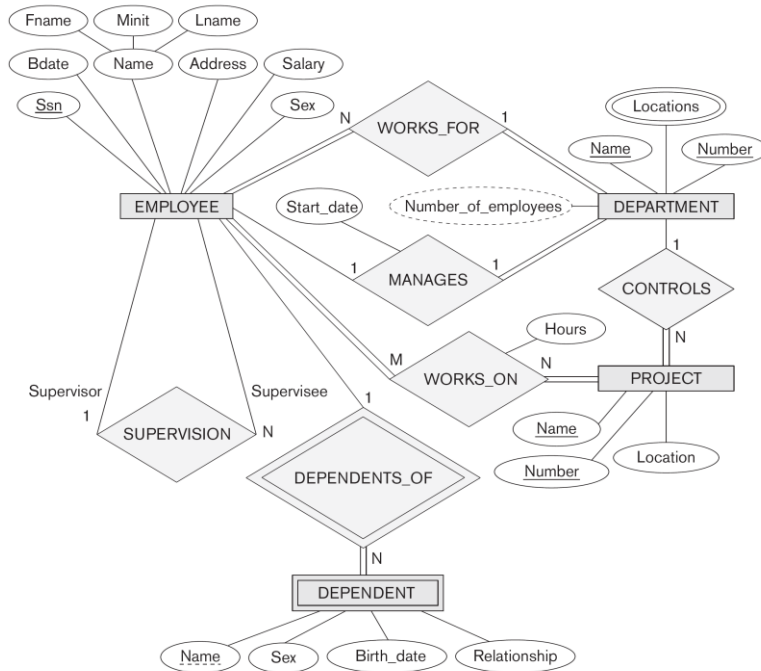




Figure 9.2
Result of mapping the
COMPANY ER schema
into a relational database
schema.



OR



Nested Queries (cont'd.)

Q4A: **SELECT** **DISTINCT** Pnumber
 FROM **PROJECT**
 WHERE Pnumber **IN**
 (**SELECT** Pnumber
 FROM **PROJECT, DEPARTMENT, EMPLOYEE**
 WHERE Dnum=Dnumber **AND**
 Mgr_ssn=Ssn **AND** Lname='Smith')

 OR
 Pnumber **IN**
 (**SELECT** Pno
 FROM **WORKS_ON, EMPLOYEE**
 WHERE Essn=Ssn **AND** Lname='Smith');

Nested Queries (cont'd.)

- Use tuples of values in comparisons
 - Place them within parentheses


```
SELECT    DISTINCT Essn
FROM      WORKS_ON
WHERE     (Pno, Hours) IN ( SELECT    Pno, Hours
                           FROM      WORKS_ON
                           WHERE     Essn='123456789' );
```

This query will select the Essns of all employees who their work pattern (project, hours) is similar to employee 'John Smith' (whose Ssn = '123456789')

Nested Queries (cont'd.)

- Use other comparison operators to compare a single value v
 - = ANY (or = SOME) operator
 - Returns TRUE if the value v is equal to some value in the set V and is hence equivalent to IN
 - Other operators that can be combined with ANY (or SOME): >, >=, <, <=, and <>

```
SELECT  Lname, Fname
FROM    EMPLOYEE
WHERE   Salary > ALL ( SELECT  Salary
                        FROM    EMPLOYEE
                        WHERE   Dno=5 );
```



Nested Queries (cont'd.)

- Avoid potential errors and ambiguities
 - Create tuple variables (aliases) for all tables referenced in SQL query

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:  SELECT    E.Fname, E.Lname
      FROM      EMPLOYEE AS E
      WHERE     E.Ssn IN ( SELECT    Essn
                          FROM      DEPENDENT AS D
                          WHERE     E.Fname=D.Dependent_name
                          AND E.Sex=D.Sex );
```

Correlated Nested Queries

- Whenever a condition in the WHERE clause of a nested query references some attribute of a relation declared in the outer query, the two queries are said to be **correlated**.
- **Correlated** nested query
 - Evaluated once for each tuple in the outer query

The EXISTS and UNIQUE Functions in SQL

- EXISTS function
 - Check whether the result of a correlated nested query is empty or not
- EXISTS and NOT EXISTS
 - Typically used in conjunction with a correlated nested query
- SQL function UNIQUE (Q)
 - Returns TRUE if there are no duplicate tuples in the result of query Q

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

Formulate Query 16 in an alternative form that uses EXISTS as in Q16B:

```
Q16B:  SELECT    E.Fname, E.Lname
        FROM      EMPLOYEE AS E
        WHERE     EXISTS ( SELECT    *
                           FROM      DEPENDENT AS D
                           WHERE     E.Ssn=D.Essn AND E.Sex=D.Sex
                                   AND E.Fname=D.Dependent_name);
```

Explicit Sets and Renaming of Attributes in SQL


- Can use explicit set of values in WHERE clause
- Use qualifier AS followed by desired new name
 - Rename any attribute that appears in the result of a query

```
Q8A:  SELECT  E.Lname AS Employee_name, S.Lname AS Supervisor_name
        FROM    EMPLOYEE AS E, EMPLOYEE AS S
        WHERE   E.Super_ssn=S.Ssn;
```

Joined Tables in SQL and Outer Joins

- **Joined table**
 - Permits users to specify a table resulting from a join operation in the FROM clause of a query
- The FROM clause in Q1A
 - Contains a single joined table

Q1A: SELECT Fname, Lname, Address
 FROM (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
 WHERE Dname='Research';



Joined Tables in SQL and Outer Joins (cont'd.)

- Specify different types of join
 - NATURAL JOIN
 - Various types of OUTER JOIN
- NATURAL JOIN on two relations R and S
 - No join condition specified
 - Implicit EQUIJOIN condition for each pair of attributes with same name from R and S

Natural Join Example

R

A	B
X	Y
X	Z
Y	Z
Z	V

S

B	C
Z	U
V	W
Z	V

A	B	C
X	Z	U
X	Z	V
Y	Z	U
Y	Z	V
Z	V	W

Natural Join

Loan

<u>Loan_no</u>	Branch_name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

Borrower

Customer_name	Loan_no
Jones	L-170
Smith	L-230
Hayes	L-155

Loan \bowtie Borrower

<u>Loan_no</u>	Branch_name	amount	Customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

Joined Tables in SQL and Outer Joins (cont'd.)

- **Inner join**

- Most frequently used type of join
- Referred to as an EQUIJOIN.
- query compares each row of table1 with each row of table2 to find all pairs of rows which satisfy the join-predicate.

Inner Join Example

EMPLOYEES Table

EMPNO	LASTNAME	WORKDEPT
001	JAGGER	A01
002	RICHARDS	M01
003	WOOD	M01
004	WATTS	C01
005	WYMAN	-
006	JONES	S01

DEPARTMENT Table

DEPTNO	DEPTNAME
A01	ADMINISTRATIVE
E01	ENGINEERING
M01	MANUFACTURING
S01	MARKETING
S02	SALES
C01	CUSTOMER SUPPORT

Inner Join

```
SELECT lastname, deptname  
FROM employees e INNER JOIN department d  
ON e.workdept = d.deptno
```

Result Data Set

LASTNAME	DEPTNAME
JAGGER	ADMINISTRATIVE
RICHARDS	MANUFACTURING
WOOD	MANUFACTURING
WATTS	CUSTOMER SUPPORT
JONES	MARKETING

Joined Tables in SQL and Outer Joins (cont'd.)

- **LEFT OUTER JOIN**
 - Every tuple in left table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of right table
- **RIGHT OUTER JOIN**
 - Every tuple in right table must appear in result
 - If no matching tuple
 - Padded with NULL values for the attributes of left table
- **FULL OUTER JOIN**

Left Outer join

Loan

<u>Loan_no</u>	Branch_name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

Borrower

Customer_name	Loan_no
Jones	L-170
Smith	L-230
Hayes	L-155

Loan \bowtie Borrower

Loan_no	Branch_name	amount	Customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	Null

Right Outer Join

Loan

<u>Loan_no</u>	Branch_name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

Borrower

Customer_name	Loan_no
Jones	L-170
Smith	L-230
Hayes	L-155

Loan \bowtie Borrower

Loan_no	Branch_name	amount	Customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	Null	Null	Hayes

Full Outer Join

Loan

<u>Loan_no</u>	Branch_name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700


Borrower

Customer_name	Loan_no
Jones	L-170
Smith	L-230
Hayes	L-155


Loan \times Borrower

Loan_no	Branch_name	amount	Customer_name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	Null
L-155	Null	Null	Hayes


(natural join)




(left outer join)



(right outer join)



(full outer join)



Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- **Grouping**
 - Create subgroups of tuples before summarizing
- Built-in aggregate functions
 - **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- Functions can be used in the `SELECT` clause or in a `HAVING` clause

Aggregate Functions in SQL

Query 20. Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.


Q20: **SELECT** **SUM** (Salary), **MAX** (Salary), **MIN** (Salary), **AVG** (Salary)
 FROM (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
 WHERE Dname='Research';

EMPLOYEE

Frame	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------



Aggregate Functions in SQL

Queries 21 and 22. Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

Q21: **SELECT** **COUNT (*)**
 FROM **EMPLOYEE;**

Q22: **SELECT** **COUNT (*)**
 FROM **EMPLOYEE, DEPARTMENT**
 WHERE **DNO=DNUMBER AND DNAME='Research';**

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------



Grouping: The GROUP BY and HAVING Clauses

- **Partition** relation into subsets of tuples
 - Based on **grouping attribute(s)**
 - Apply function to each such group independently
- **GROUP BY** clause
 - Specifies grouping attributes
- If NULLs exist in grouping attribute
 - Separate group created for all tuples with a NULL value in grouping attribute

“Group By” Example → 1 attribute

Employee

SSN	FName	other attributes	Sex	DNO	Salary
111-22-3333	John	•••••	M	4	40000
123-45-6789	Mary	•••••	F	5	50000
987-82-9823	James	•••••	M	5	60000
982-71-9927	Jake	•••••	M	4	50000

Group by DNO

111-22-3333	John	•••••	M	4	40000
982-71-9927	Jake	•••••	M	4	50000

123-45-6789	Mary	•••••	F	5	50000
987-82-9823	James	•••••	M	5	60000


“Group By” Example → more than 1 attribute

Employee

SSN	FName	other attributes	Sex	DNO	Salary
111-22-3333	John	•••••	M	4	40000
123-45-6789	Mary	•••••	F	5	50000
987-82-9823	James	•••••	M	5	60000
982-71-9927	Jake	•••••	M	4	50000



Group by (DNO, Sex)



Group By and Aggregate Function

Employees

DEPARTMENT_ID	SALARY		DEPARTMENT_ID	SUM(SALARY)
10	4400	4400	10	4400
20	13000	19000	20	19000
20	6000		30	24900
30	11000	24900	40	6500
30	3100		50	156400
30	2900			
30	2800			
30	2600			
30	2500			
40	6500	6500		
50	8000	156400		
50	8200			
50	7900			
50	6500			
50	5800			

Sum of Salary in Employees table for each Department

Group By and Aggregate Function

FoodChart

date	food	sold
06/05/13	pizza	349
06/06/13	hotdog	500
06/06/13	pizza	70

SELECT food, sum(sold) as totalSold
FROM Foodchart
group by food;

food	totalSold
hotdog	500
pizza	419

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

- **HAVING** clause
 - Provides a condition on the summary information

```
SELECT CUST_CITY,COUNT(*) FROM CUSTOMERS  
GROUP BY CUST_CITY;
```

CUST_CITY	Count(*)
DELHI	3
LUCKNOW	2
CHENNAI	1
BANGALORE	3
MUMBAI	1

```
SELECT CUST_CITY,COUNT(*) FROM CUSTOMERS  
GROUP BY CUST_CITY HAVING COUNT(*) > 2;
```

CUST_CITY	Count(*)
DELHI	3
BANGALORE	3

Discussion and Summary of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

Quiz 1

- Consider the following relational schema. An employee can work in more than one department; the *pct time* field of the Works relation shows the percentage of time that a given employee works in a given department.

`Emp(eid: integer, ename: string, age: integer, salary: real)`

`Works(eid: integer, did: integer, pct_time: integer)`

`Dept(did: integer, dname: string, budget: real, managerid: integer)`

Quiz 1

- Write the following queries in SQL:
 - Print the names and ages of each employee who works in both the Hardware department and the Software department.
 - For each department with more than 20 full-time-equivalent employees (i.e., where the part-time and full-time employees add up to at least that many full-time employees), print the *did* together with the number of employees that work in that department.

Quiz 1

- Print the name of each employee whose salary exceeds the budget of all of the departments that he or she works in.
- Find the *enames* of managers who manage only departments with budgets larger than \$1 million, but at least one department with budget less than \$5 million.

SQL Advanced

- More Complex SQL Retrieval Queries
- Specifying Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Change Statements in SQL



Specifying Constraints as Assertions and Actions as Triggers

- **CREATE ASSERTION**

- Specify additional types of constraints outside scope of built-in relational model constraints

- **CREATE TRIGGER**

- Specify automatic actions that database system will perform when certain events and conditions occur

Specifying General Constraints as Assertions in SQL

- CREATE ASSERTION

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT *
                     FROM   EMPLOYEE E, EMPLOYEE M,
                          DEPARTMENT D
                     WHERE  E.Salary>M.Salary
                          AND E.Dno=D.Dnumber
                          AND D.Mgr_ssn=M.Ssn ) );
```

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------



Introduction to Triggers in SQL

- `CREATE TRIGGER` statement
 - Used to monitor the database
- Typical trigger has three components:
 - **Event(s)**
 - **Condition**
 - **Action**

Triggers in SQL

- Check whenever an employee's salary is greater than the salary of his or her direct supervisor in the COMPANY database
- Event:
 - Inserting a new employee record,
 - Changing an employee's salary,
 - Changing an employee's supervisor

Triggers in SQL

- Action:
 - Execute the stored procedure INFORM_SUPERVISOR.
- Following trigger in oracle:

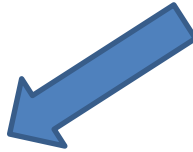
```
R5: CREATE TRIGGER SALARY_VIOLATION
      BEFORE INSERT OR UPDATE OF SALARY, SUPERVISOR_SSN
      ON EMPLOYEE
      FOR EACH ROW
      WHEN ( NEW.SALARY > ( SELECT SALARY FROM EMPLOYEE
                           WHERE SSN = NEW.SUPERVISOR_SSN ) )
      INFORM_SUPERVISOR(NEW.Supervisor_ssn,
                        NEW.Ssn );
```

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno



SQL Advanced

- More Complex SQL Retrieval Queries
- Specifying Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL 
- Schema Change Statements in SQL

Views (Virtual Tables) in SQL

- Concept of a view in SQL
 - Single table derived from other tables
 - Considered to be a virtual table

Specification of Views in SQL

- **CREATE VIEW** command
 - Give table name, list of attribute names, and a query to specify the contents of the view

V1: **CREATE VIEW** **WORKS_ON1**
AS SELECT Fname, Lname, Pname, Hours
FROM EMPLOYEE, PROJECT, WORKS_ON
WHERE Ssn=Essn **AND** Pno=Pnumber;

EMPLOYEE


Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

PROJECT

Pname	<u>Pnumber</u>	<u>Plocation</u>	Dnum
-------	----------------	------------------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------



```

V2:      CREATE VIEW  DEPT_INFO(Dept_name, No_of_emps, Total_sal)
        AS SELECT    Dname, COUNT (*), SUM (Salary)
                FROM    DEPARTMENT, EMPLOYEE
                WHERE    Dnumber=Dno
                GROUP BY Dname;

```

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Edate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------



Specification of Views in SQL (cont'd.)

- Specify SQL queries on a view
- View always up-to-date
 - Responsibility of the DBMS and not the user
- **DROP VIEW** command
 - Dispose of a view

View Implementation, View Update, and Inline Views

- Complex problem of efficiently implementing a view for querying
- **Query modification** approach
 - Modify view query into a query on underlying base tables
 - Disadvantage: inefficient for views defined via complex queries that are time-consuming to execute

View Implementation

- **View materialization approach**
 - Physically create a temporary view table when the view is first queried
 - Keep that table on the assumption that other queries on the view will follow
 - Requires efficient strategy for automatically updating the view table when the base tables are updated

View Implementation (cont'd.)


- **Incremental update strategies**
 - DBMS determines what new tuples must be inserted, deleted, or modified in a materialized view table

View Update and Inline Views

- Update on a view defined on a single table without any aggregate functions
 - Can be mapped to an update on underlying base table
- View involving joins
 - Often not possible for DBMS to determine which of the updates is intended

SQL Advanced

- More Complex SQL Retrieval Queries
- Specifying Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Change Statements in SQL



Schema Change Statements in SQL

- **Schema evolution commands**
 - Can be done while the database is operational
 - Does not require recompilation of the database schema

The DROP Command

- DROP command
 - Used to drop named schema elements, such as tables, domains, or constraint
- Drop behavior options:
 - CASCADE and RESTRICT
- Example:
 - DROP SCHEMA COMPANY CASCADE;
 - Automatically drop objects (tables, functions, etc.) that are contained in the schema.
-

The ALTER Command

- **Alter table actions** include:
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints
- **Example:**
 - `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12) ;`
- **To drop a column**
 - Choose either `CASCADE` or `RESTRICT`

- When you remove a column from a table, will automatically remove all of the indexes and constraints that involved the dropped column.
- If the column that you want to remove is used in other database objects such views, triggers, stored procedures, etc., you cannot drop the column because other objects are depending on it.
 - In this case, you need to add the CASCADE option to the DROP COLUMN clause to drop the column and all of its dependent objects:

The ALTER Command (cont'd.)

- Change constraints specified on a table
 - Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```


Summary

- Complex SQL:
 - Nested queries, joined tables, outer joins, aggregate functions, grouping
- CREATE ASSERTION and CREATE TRIGGER
- Views
 - Virtual or derived tables

Quiz 2

- Consider the following relational schema and briefly answer the questions that follow:

`Emp(eid: integer, ename: string, age: integer, salary: real)`

`Works(eid: integer, did: integer, pct_time: integer)`

`Dept(did: integer, budget: real, managerid: integer)`

- Next Page->

Quiz 2

- 1. Define a table constraint on Emp that will ensure that every employee makes at least \$10,000.
- 2. Define a table constraint on Dept that will ensure that all managers have *age* > 30.
- 3. Define an assertion on Dept that will ensure that all managers have *age* > 30. Compare this assertion with the equivalent table constraint. Explain which is better.

- CREATE TABLE Emp (eid INTEGER, ename CHAR(10), age INTEGER , salary REAL, PRIMARY KEY (eid), CHECK (*salary* >= 10000))